

# ResNet10

December 10, 2025

```
[1]: import os
import glob
import numpy as np
import cv2
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, roc_auc_score, f1_score, ▾
    confusion_matrix, accuracy_score
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: np.random.seed(42)
tf.random.set_seed(42)

# Configuration
POSITIVE_DIR = "/content/DDS1/PP/IMAGE"
NEGATIVE_DIR = "/content/DDS1/NP/IMAGE"

# Global Params
IMG_WIDTH = 40
IMG_HEIGHT = 40
IMG_CHANNELS = 3
IMG_SHAPE = (IMG_HEIGHT, IMG_WIDTH, IMG_CHANNELS)

TEST_SPLIT_SIZE = 0.2
VALIDATION_SPLIT = 0.2
EPOCHS = 30
BATCH_SIZE = 32
DROPOUT_RATE = 0.2 # Best dropout
```

```
[3]: def load_and_preprocess_images(positive_dir, negative_dir):
    images = []
    labels = []

    print(f"Loading images from: {positive_dir} (Label 1)")
```

```

pos_files = glob.glob(os.path.join(positive_dir, "*.bmp"))
print(f"Found {len(pos_files)} positive images.")

for img_path in pos_files:
    try:
        img = cv2.imread(img_path)
        if img is None: continue
        img_resized = cv2.resize(img, (IMG_WIDTH, IMG_HEIGHT))
        if len(img_resized.shape) == 2 or img_resized.shape[2] == 1:
            img_resized = cv2.cvtColor(img_resized, cv2.COLOR_GRAY2RGB)
        images.append(img_resized)
        labels.append(1)
    except Exception as e: pass

print(f"Loading images from: {negative_dir} (Label 0)")
neg_files = glob.glob(os.path.join(negative_dir, "*.bmp"))
print(f"Found {len(neg_files)} negative images.")

for img_path in neg_files:
    try:
        img = cv2.imread(img_path)
        if img is None: continue
        img_resized = cv2.resize(img, (IMG_WIDTH, IMG_HEIGHT))
        if len(img_resized.shape) == 2 or img_resized.shape[2] == 1:
            img_resized = cv2.cvtColor(img_resized, cv2.COLOR_GRAY2RGB)
        images.append(img_resized)
        labels.append(0)
    except Exception as e: pass

print("Converting to NumPy arrays and normalizing...")
images_np = np.array(images, dtype="float32") / 255.0
labels_np = np.array(labels, dtype="int32")
return images_np, labels_np

```

```
[4]: def residual_block(x, filters, stride=1):
    shortcut = x
    x = layers.Conv2D(filters, (3, 3), strides=stride, padding='same')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('relu')(x)
    x = layers.Conv2D(filters, (3, 3), strides=1, padding='same')(x)
    x = layers.BatchNormalization()(x)

    if stride != 1 or shortcut.shape[-1] != filters:
        shortcut = layers.Conv2D(filters, (1, 1), strides=stride, padding='same')(shortcut)
        shortcut = layers.BatchNormalization()(shortcut)
```

```

x = layers.Add()([x, shortcut])
x = layers.Activation('relu')(x)
return x

def build_winning_model(input_shape, dropout_rate):
    inputs = keras.Input(shape=input_shape, name="input_image")

    # Initial Conv
    x = layers.Conv2D(64, (3, 3), strides=1, padding='same')(inputs)
    x = layers.BatchNormalization()(x)
    x = layers.Activation('relu')(x)

    # ResNet-10 Blocks
    x = residual_block(x, 64, stride=1)
    x = residual_block(x, 64, stride=2)
    x = residual_block(x, 128, stride=2)
    x = residual_block(x, 256, stride=2)

    # Head
    x = layers.GlobalAveragePooling2D()(x)
    feature_output = layers.Dense(128, activation='relu', ↴
        name='feature_layer')(x)
    x = layers.Dropout(dropout_rate)(feature_output)
    classifier_output = layers.Dense(1, activation='sigmoid', ↴
        name='classifier_head')(x)

    model = keras.Model(inputs=inputs, outputs=classifier_output, ↴
        name="ResNet10_Hero")
    model.compile(optimizer='adam', loss='binary_crossentropy', ↴
        metrics=['accuracy'])
    return model

def plot_confusion_matrix(y_true, y_pred, title):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=['Negative', 'Positive'],
                yticklabels=['Negative', 'Positive'])
    plt.title(title)
    plt.ylabel('Actual')
    plt.xlabel('Predicted')
    plt.show()

```

```
[5]: def run_resnet():
    print(f"--- STARTING RUN (ResNet-10, Dropout {DROPOUT_RATE}) ---")
```

```
# Load Data
```

```

try:
    all_images, all_labels = load_and_preprocess_images(POSITIVE_DIR, □
NEGATIVE_DIR)
except Exception as e:
    print(f"Error loading images: {e}")
    return

if len(all_images) == 0: return

# Split
X_train, X_test, y_train, y_test = train_test_split(
    all_images, all_labels, test_size=TEST_SPLIT_SIZE, stratify=all_labels, □
random_state=42
)

# Build Model
model = build_winning_model(IMG_SHAPE, DROPOUT_RATE)

# Class Weights (1:4)
class_weights_dict = {0: 1.0, 1: 4.0}

# Early Stopping
early_stop = keras.callbacks.EarlyStopping(
    monitor='val_loss',
    patience=5, # Give it 5 epochs to improve
    restore_best_weights=True
)

# Train
print("\nTraining ResNet Model...")
history = model.fit(
    X_train, y_train,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    validation_split=VALIDATION_SPLIT,
    class_weight=class_weights_dict,
    callbacks=[early_stop],
    verbose=1
)

# Evaluate
print("\nEvaluating on Test Set...")
preds_proba = model.predict(X_test).ravel()
preds_class = (preds_proba > 0.5).astype(int)

# Metrics
final_auc = roc_auc_score(y_test, preds_proba)

```

```

final_f1 = f1_score(y_test, preds_class)
final_acc = accuracy_score(y_test, preds_class)

print("\n==== FINAL RESNET MODEL RESULTS ===")
print(f"AUC Score: {final_auc:.5f}")
print(f"F1 Score: {final_f1:.5f}")
print(f"Accuracy: {final_acc:.5f}")

print("\nClassification Report:")
print(classification_report(y_test, preds_class, target_names=['Negative', ↴'Positive']))

plot_confusion_matrix(y_test, preds_class, "ResNet-10 Confusion Matrix")

```

```
[6]: if __name__ == "__main__":
    run_resnet()
```

```

--- STARTING RUN (ResNet-10, Dropout 0.2) ---
Loading images from: /content/DDS1/PP/IMAGE (Label 1)
Found 2636 positive images.
Loading images from: /content/DDS1/NP/IMAGE (Label 0)
Found 28848 negative images.
Converting to NumPy arrays and normalizing...

Training ResNet Model...
Epoch 1/30
630/630          37s 30ms/step -
accuracy: 0.9380 - loss: 0.3002 - val_accuracy: 0.9843 - val_loss: 0.0466
Epoch 2/30
630/630          5s 8ms/step -
accuracy: 0.9846 - loss: 0.0803 - val_accuracy: 0.9369 - val_loss: 0.1814
Epoch 3/30
630/630          5s 8ms/step -
accuracy: 0.9890 - loss: 0.0577 - val_accuracy: 0.9833 - val_loss: 0.0487
Epoch 4/30
630/630          5s 8ms/step -
accuracy: 0.9920 - loss: 0.0467 - val_accuracy: 0.9764 - val_loss: 0.0753
Epoch 5/30
630/630          5s 8ms/step -
accuracy: 0.9910 - loss: 0.0443 - val_accuracy: 0.9857 - val_loss: 0.0466
Epoch 6/30
630/630          5s 8ms/step -
accuracy: 0.9925 - loss: 0.0362 - val_accuracy: 0.9869 - val_loss: 0.0390
Epoch 7/30
630/630          5s 8ms/step -
accuracy: 0.9915 - loss: 0.0421 - val_accuracy: 0.9925 - val_loss: 0.0236
Epoch 8/30
630/630          5s 8ms/step -

```

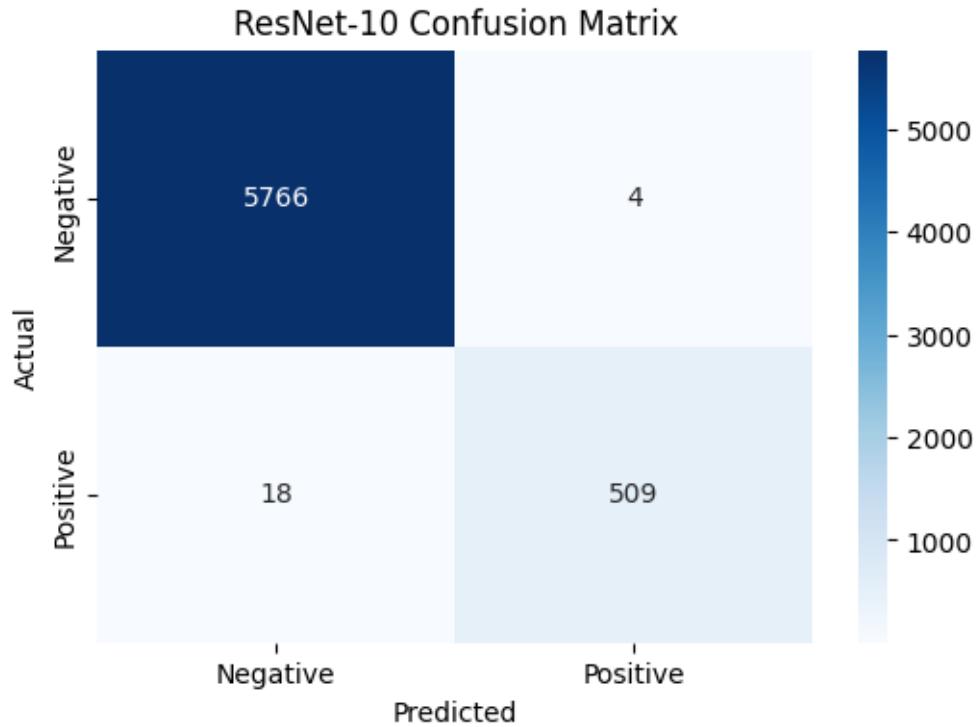
```
accuracy: 0.9942 - loss: 0.0294 - val_accuracy: 0.9891 - val_loss: 0.0337
Epoch 9/30
630/630          5s 8ms/step -
accuracy: 0.9954 - loss: 0.0283 - val_accuracy: 0.9927 - val_loss: 0.0244
Epoch 10/30
630/630          5s 8ms/step -
accuracy: 0.9954 - loss: 0.0254 - val_accuracy: 0.9861 - val_loss: 0.0503
Epoch 11/30
630/630          5s 8ms/step -
accuracy: 0.9946 - loss: 0.0279 - val_accuracy: 0.9946 - val_loss: 0.0173
Epoch 12/30
630/630          5s 8ms/step -
accuracy: 0.9944 - loss: 0.0275 - val_accuracy: 0.9522 - val_loss: 0.1417
Epoch 13/30
630/630          5s 8ms/step -
accuracy: 0.9934 - loss: 0.0343 - val_accuracy: 0.9764 - val_loss: 0.0606
Epoch 14/30
630/630          5s 8ms/step -
accuracy: 0.9962 - loss: 0.0191 - val_accuracy: 0.9865 - val_loss: 0.0376
Epoch 15/30
630/630          5s 9ms/step -
accuracy: 0.9964 - loss: 0.0189 - val_accuracy: 0.9942 - val_loss: 0.0227
Epoch 16/30
630/630          5s 8ms/step -
accuracy: 0.9952 - loss: 0.0210 - val_accuracy: 0.9847 - val_loss: 0.0556
```

```
Evaluating on Test Set...
197/197          5s 16ms/step
```

```
==== FINAL RESNET MODEL RESULTS ====
AUC Score: 0.99966
F1 Score: 0.97885
Accuracy: 0.99651
```

#### Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Negative     | 1.00      | 1.00   | 1.00     | 5770    |
| Positive     | 0.99      | 0.97   | 0.98     | 527     |
| accuracy     |           |        | 1.00     | 6297    |
| macro avg    | 0.99      | 0.98   | 0.99     | 6297    |
| weighted avg | 1.00      | 1.00   | 1.00     | 6297    |



```
[11]: from IPython.display import display

hero_data = {
    'Metric': [
        'AUC Score',
        'F1-Score (Positive)',
        'Recall (Sensitivity)',
        'Precision (Positive)',
        'Overall Accuracy'
    ],
    'Value': [
        0.9997, # AUC
        0.9789, # F1 Score
        0.9700, # Recall for Positive Class (from report)
        0.9900, # Precision for Positive Class (from report)
        0.9965 # Accuracy
    ]
}

# Create DataFrame
df_hero = pd.DataFrame(hero_data)

# Transpose it so it looks like a summary row
```

```

hero_data_horizontal = {
    'Model Configuration': ['ResNet-10 (Dropout 0.2)'],
    'AUC': [0.9997],
    'F1-Score': [0.9789],
    'Recall': [0.9700],
    'Precision': [0.9900],
    'Accuracy': [0.9965]
}

df_hero_final = pd.DataFrame(hero_data_horizontal)

# Display table
print("--- Table II: Final Validated Model Performance ---")
display(df_hero_final.style.hide(axis='index') # Hides the '0' index number
        .background_gradient(cmap='Greens', subset=['AUC', 'F1-Score', 'Recall'])
        .format("{:.4f}", subset=['AUC', 'F1-Score', 'Recall', 'Precision', 'Accuracy']))

```

--- Table II: Final Validated Model Performance ---

<pandas.io.formats.style.Styler at 0x7ed6e7d2e5a0>